

MACHINE LEARNING

(Predict the Forest Fires)

Summer Internship Report Submitted in partial fulfillment

of the requirement for undergraduate degree of

Bachelor of Technology

In

Computer Science Engineering

By

Gajjala Maansi

221710310017

Under the Guidance of

Assistant Professor



Department Of Computer Science Engineering

GITAM School of Technology

GITAM (Deemed to be University)

Hyderabad-502329

June 2020

DECLARATION

I submit this industrial training work entitled “**PREDICT THE FOREST FIRES**” to GITAM (Deemed To Be University), Hyderabad in partial fulfillment of the requirements for the award of the degree of “**Bachelor of Technology**” in “**Computer Science Engineering**”. I declare that it was carried out independently by me under the guidance of _____, Asst. Professor, GITAM (Deemed To Be University), Hyderabad, India.

The results embodied in this report have not been submitted to any other University or Institute for the award of any degree or diploma.

Place: HYDERABAD

Gajjala Maansi

Date:

221710310017



CERTIFICATE

This is to certify that the Industrial Training Report entitled **“PREDICT THE FOREST FIRES”** is being submitted by GAJJALA MAANSI (221710310017) in partial fulfillment of the requirement for the award of **Bachelor of Technology in Computer Science Engineering** at GITAM (Deemed To Be University), Hyderabad during the academic year 2020-21

It is faithful record work carried out by her at the **Computer Science Engineering Department**, GITAM University Hyderabad Campus under my guidance and supervision.

Assistant Professor
Department of CSE

Dr. S. Phani Kumar
Professor and HOD
Department of CSE

ACKNOWLEDGEMENT

Apart from my effort, the success of this internship largely depends on the encouragement and guidance of many others. I take this opportunity to express my gratitude to the people who have helped me in the successful completion of this internship.

I would like to thank respected **Dr. N. Siva Prasad**, Pro Vice Chancellor, GITAM Hyderabad and **Dr. CH. Sanjay**, Principal, GITAM Hyderabad

I would like to thank respected **Dr. S. Phani Kumar**, Head of the Department of Computer Science Engineering for giving me such a wonderful opportunity to expand my knowledge for my own branch and giving me guidelines to present a internship report. It helped me a lot to realize of what we study for.

I would like to thank the respected faculties who helped me to make this internship a successful accomplishment.

I would also like to thank my friends who helped me to make my work more organized and well-stacked till the end.

Gajjala Maansi

221710310017

ABSTRACT

Forest fires are a major environmental issue, creating economical and ecological damage while endangering human lives. Fast detection is a key element for controlling such phenomenon. To achieve this, one alternative is to use automatic tools based on local sensors, such as provided by meteorological stations. In effect, meteorological conditions (e.g. temperature, wind) are known to influence forest fires and several fire indexes, such as the forest Fire Weather Index (FWI), use such data.

In this work, we explore a Data Mining (DM) approach and Linear Regression to predict the burned area of forest fires. Five different DM techniques, e.g. Random Forests and Neural Networks and four distinct feature selection setups (using spatial, temporal, FWI components and weather attributes), were tested on recent real-world data collected from the northeast region of Portugal. It is capable of predicting the burned area of small fires, which are more frequent. Such knowledge is particularly useful for improving firefighting resource management (e.g. prioritizing targets for air tankers and ground crews).

Table of Contents:

CHAPTER 1:MACHINE LEARNING.....	1
1.1 INTRODUCTION.....	1
1.2 IMPORTANCE OF MACHINE LEARNING.....	1
1.3 USES OF MACHINE LEARNING.....	2
1.4 TYPES OF LEARNING ALGORITHMS.....	3
1.4.1 Supervised Learning.....	3
1.4.2 Unsupervised Learning.....	4
1.4.3 Semi Supervised Learning.....	5
1.5 RELATION BETWEEN DATA MINING,MACHINE LEARNING AND DEEP LEARNING.....	5
CHAPTER 2:PYTHON.....	7
2.1 INTRODUCTION TO PYTHON.....	7
2.2 HISTORY OF PYTHON.....	7
2.3 FEATURES OF PYTHON.....	8
2.4 HOW TO SETUP PYTHON.....	8
2.4.1 Installation(using python IDLE).....	8
2.4.2 Installation(using Anaconda).....	9
2.5 PYTHON VARIABLE TYPES.....	10
2.5.1 Python Numbers.....	12
2.5.2 Python Strings.....	12
2.5.3 Python Lists.....	12
2.5.4 Python Tuples.....	13
2.5.5 Python Dictionary.....	13
2.6 PYTHON FUNCTION.....	13
2.6.1 Defining a Function.....	13
2.6.2 Calling a Function.....	15
2.7 PYTHON USING OOP's CONCEPTS.....	15
2.7.1 Class.....	15
2.7.2 __init__ method in class.....	16
CHAPTER 3:CASE STUDY.....	17
3.1 PROBLEM STATEMENT.....	17
3.2 DATA SET.....	17
3.3 OBJECTIVE OF THE CASE STUDY.....	18
CHAPTER 4:MODEL BUILDING.....	19
4.1 PREPROCESSING OF THE DATA.....	19
4.1.1 Getting the Data Set.....	19
4.1.2 Importing the Libraries.....	19
4.1.3 Importing the Data-Set.....	19
4.1.4 Handling the Missing values.....	20
4.1.5 Data Visualization.....	21
4.1.6 Categorical Data.....	27

4.2 TRAINING THE MODEL.....	33
4.3 EVALUATING THE CASE STUDY.....	34
4.3.1 Building the model(using splitting).....	34
CONCLUSION.....	48
REFERENCES.....	49

List of figures:

Figure 1: The process flow.....	2
Figure 2:Unsupervised Learning.....	4
Figure 3:Semi supervised Learning.....	5
Figure 4:Python download.....	9
Figure 5:Anaconda download.....	10
Figure 6:Jupyter Notebook.....	11
Figure 7:Defining a class.....	16
Figure 8:Importing libraries.....	19
Figure 9:Reading the dataset.....	20
Figure 10:Missing values.....	20
Figure 11:Scatterplots and distributions of numerical features to see how they may effect the output 'area'.....	25
Figure 12: Boxplot of how categorical column day affect the outcome.....	26
Figure 13: Boxplot of how categorical column day affect the outcome.....	26
Figure14:Categorical data of column 'month'.....	28
Figure15:Dummysset for column 'month'.....	28
Figure 16:Categorical data of column 'day'.....	29
Figure 17:Dummysset for column 'day'.....	29
Figure 18:Concatenating dummy sets to dataframe.....	30
Figure 19:Importing label encoder and one hot encoder.....	31
Figure 20:Handling categorical data of column month.....	31
Figure 21: Handling categorical data of column day.....	32
Figure 22:Importing train_test_split.....	34
Figure 23:Retrieving the input column.....	34
Figure 24:Retrieving the output column.....	35
Figure 25:Defining REC.....	35
Figure 26:Importing Random Forest Regressor and GridsearchCV...36	36
Figure 27:Parameter grid for gridsearch and rfr.....	36
Figure 28:Best parameter obtained by gridsearch.....	37
Figure 29:RMSE for rfr.....	37
Figure 30:Scatter plot for rfr.....	37
Figure 31:Histogram for prediction errors of rfr.....	38
Figure 32:REC curve for rfr.....	39
Figure 33:Importing NN packages.....	40
Figure 34:Dividing data and target.....	40
Figure 35:RMSE for NN.....	41
Figure 36:Scatterplot for NN.....	41

Figure 37:Histogram of prediction error of NN.....	42
Figure 38:REC curve for NN.....	43
Figure 39:Relative performance of rfr and NN(REC curves)....	44
Figure 40:Importing Linear Regression package.....	45
Figure 41:Predicting the output.....	46
Figure 42:Comparing y and y_pred.....	47
Figure 43:R2_score to check the model performance.....	47

CHAPTER 1

MACHINE LEARNING

INTRODUCTION:

Machine Learning(ML) is the scientific study of algorithms and statistical models that computer systems use in order to perform a specific task effectively without using explicit instructions, relying on patterns and inference instead. It is seen as a subset of Artificial Intelligence(AI).

IMPORTANCE OF MACHINE LEARNING:

Consider some of the instances where machine learning is applied: the self-driving Google car, cyber fraud detection, online recommendation engines—like friend suggestions on Facebook, Netflix showcasing the movies and shows you might like, and “more items to consider” and “get yourself a little something” on Amazon—are all examples of applied machine learning. All these examples echo the vital role machine learning has begun to take in today’s data-rich world.

Machines can aid in filtering useful pieces of information that help in major advancements, and we are already seeing how this technology is being implemented in a wide variety of industries.

With the constant evolution of the field, there has been a subsequent rise in the uses, demands, and importance of machine learning. Big data has become quite a buzzword in the last few years; that’s in part due to increased sophistication of machine learning, which helps analyze those big chunks of big data. Machine learning has also changed the way data extraction, and interpretation is done by involving automatic sets of generic methods that have replaced traditional statistical techniques.

The process flow depicted here represents how machine learning works

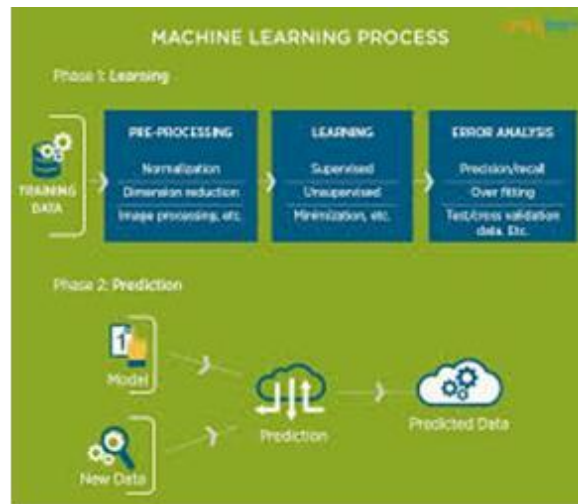


Figure 1 : The Process Flow

USES OF MACHINE LEARNING:

Earlier in this article, we mentioned some applications of machine learning. To understand the concept of machine learning better, let's consider some more examples: web search results, real-time ads on web pages and mobile devices, email spam filtering, network intrusion detection, and pattern and image recognition. All these are by-products of applying machine learning to analyze huge volumes of data

Traditionally, data analysis was always being characterized by trial and error, an approach that becomes impossible when data sets are large and heterogeneous. Machine learning comes as the solution to all this chaos by proposing clever alternatives to analyzing huge volumes of data.

By developing fast and efficient algorithms and data-driven models for real-time processing of data, machine learning can produce accurate results and analysis.

TYPES OF LEARNING ALGORITHMS:

The types of machine learning algorithms differ in their approach, the type of data they input and output, and the type of task or problem that they are intended to solve.

Supervised Learning :

When an algorithm learns from example data and associated target responses that can consist of numeric values or string labels, such as classes or tags, in order to later predict the correct response when posed with new examples comes under the category of supervised learning.

Supervised machine learning algorithms uncover insights, patterns, and relationships from a labelled training dataset – that is, a dataset that already contains a known value for the target variable for each record. Because you provide the machine learning algorithm with the correct answers for a problem during training, it is able to “learn” how the rest of the features relate to the target, enabling you to uncover insights and make predictions about future outcomes based on historical data.

Examples of Supervised Machine Learning Techniques are Regression, in which the algorithm returns a numerical target for each example, such as how much revenue will be generated from a new marketing campaign.

Classification, in which the algorithm attempts to label each example by choosing between two or more different classes. Choosing between two classes is called binary classification, such as determining whether or not someone will default on a loan. Choosing between more than two classes is referred to as multiclass classification.

Unsupervised Learning:

When an algorithm learns from plain examples without any associated response, leaving to the algorithm to determine the data patterns on its own. This type of algorithm tends to restructure the data into something else, such as new features that may represent a class or a new series of uncorrelated values. They are quite useful in providing humans with insights into the meaning of data and new useful inputs to supervised machine learning algorithms.




Figure 2 : Unsupervised Learning

Popular techniques where unsupervised learning is used also include self-organizing maps, nearest neighbor mapping, singular value decomposition, and k-means clustering. Basically, online recommendations, identification of data outliers, and segment text topics are all examples of unsupervised learning.

Semi Supervised Learning:

As the name suggests, semi-supervised learning is a bit of both supervised and unsupervised learning and uses both labeled and unlabeled data for training. In a typical scenario, the algorithm would use a small amount of labeled data with a large amount of unlabeled data.




Figure 3 : Semi Supervised Learning

RELATION BETWEEN DATA MINING,MACHINE LEARNING AND DEEP LEARNING:

Machine learning and data mining use the same algorithms and techniques as data mining, except the kinds of predictions vary. While data mining discovers previously unknown patterns and knowledge, machine learning reproduces known patterns and knowledge—and further automatically applies that information to data, decision-making, and actions.

Deep learning, on the other hand, uses advanced computing power and special

types of neural networks and applies them to large amounts of data to learn, understand, and identify complicated patterns. Automatic language translation and medical diagnoses are examples of deep learning.

CHAPTER 2

PYTHON

Basic programming language used for machine learning is : PYTHON

INTRODUCTION TO PYHTON:

- Python is a high-level, interpreted, interactive and object-oriented scripting language.
- Python is a general purpose programming language that is often applied in scripting roles
- Python is Interpreted: Python is processed at runtime by the interpreter. You do not need to compile your program before executing it. This is like PERL and PHP.
- Python is Interactive: You can sit at a Python prompt and interact with the interpreter directly to write your programs.
- Python is Object-Oriented: Python supports the Object-Oriented style or technique of programming that encapsulates code within objects.

HISTORY OF PYTHON:

- Python was developed by GUIDO VAN ROSSUM in early 1990's
- Its latest version is 3.7 , it is generally called as python3

FEATURES OF PYTHON:

- Easy-to-learn: Python has few keywords, simple structure, and a clearly defined syntax, This allows the student to pick up the language quickly.
- Easy-to-read: Python code is more clearly defined and visible to the eyes.
- Easy-to-maintain: Python's source code is fairly easy-to-maintaining.
- A broad standard library: Python's bulk of the library is very portable and cross-platform compatible on UNIX, Windows, and Macintosh.
- Portable: Python can run on a wide variety of hardware platforms and has the same interface on all platforms.
- Extendable: You can add low-level modules to the Python interpreter. These modules enable programmers to add to or customize their tools to be more efficient.
- Databases: Python provides interfaces to all major commercial databases.
- GUI Programming: Python supports GUI applications that can be created and ported to many system calls, libraries and windows systems, such as Windows MFC, Macintosh, and the X Window system of Unix.

HOW TO SETUP PYTHON:

- Python is available on a wide variety of platforms including Linux and Mac OS X. Let's understand how to set up our Python environment.
- The most up-to-date and current source code, binaries, documentation, news, etc., is available on the official website of Python.

Installation(using python IDLE):

- Installing python is generally easy, and nowadays many Linux and Mac OS distributions include a recent python.
- Download python from www.python.org
- When the download is completed, double click the file and follow the instructions to install it.
- When python is installed, a program called IDLE is also installed along with it. It provides a graphical user interface to work with python.



Figure 4 : Python download

Installation(using Anaconda):

- Python programs are also executed using Anaconda.
- Anaconda is a free open source distribution of python for large scale data processing, predictive analytics and scientific computing.
- Conda is a package manager quickly installs and manages packages.

- In WINDOWS:
- In windows
 - Step 1: Open Anaconda.com/downloads in web browser.
 - Step 2: Download python 3.4 version for (32-bitgraphic installer/64 -bit graphic installer)
 - Step 3: select installation type(all users)
 - Step 4: Select path(i.e. add anaconda to path & register anaconda as default python 3.4) next click install and next click finish
 - Step 5: Open jupyter notebook (it opens in default browser)




Figure 5 : Anaconda download



Figure 6 : Jupyter notebook

PYTHON VARIABLE TYPES:

- Variables are nothing but reserved memory locations to store values. This means that when you create a variable you reserve some space in memory.
- Variables are nothing but reserved memory locations to store values.
- Based on the data type of a variable, the interpreter allocates memory and decides what can be stored in the reserved memory.
- Python variables do not need explicit declaration to reserve memory space. The declaration happens automatically when you assign a value to a variable.
- Python has various standard data types that are used to define the operations possible on them and the storage method for each of them.
- Python has five standard data types –
 - Numbers
 - Strings
 - Lists

- Tuples
- Dictionary

Python Numbers:

- Number data types store numeric values. Number objects are created when you assign a value to them.
- Python supports four different numerical types – int (signed integers) long (long integers, they can also be represented in octal and hexadecimal) float (floating point real values) complex (complex numbers).

Python Strings:

- Strings in Python are identified as a contiguous set of characters represented in the quotation marks.
- Python allows for either pairs of single or double quotes.
- Subsets of strings can be taken using the slice operator ([] and [:]) with indexes starting at 0 in the beginning of the string and working their way from -1 at the end.
- The plus (+) sign is the string concatenation operator and the asterisk (*) is the repetition operator.

Python Lists:

- Lists are the most versatile of Python's compound data types.
- A list contains items separated by commas and enclosed within square brackets

([]).

- To some extent, lists are similar to arrays in C. One difference between them is that all the items belonging to a list can be of different data type.
- The values stored in a list can be accessed using the slice operator ([] and [:]) with indexes starting at 0 in the beginning of the list and working their way to end -1.
- The plus (+) sign is the list concatenation operator, and the asterisk (*) is the repetition operator.

Python Tuples:

- A tuple is another sequence data type that is similar to the list.
- A tuple consists of a number of values separated by commas. Unlike lists, however, tuples are enclosed within parentheses.
- The main differences between lists and tuples are: Lists are enclosed in brackets ([]) and their elements and size can be changed, while tuples are enclosed in parentheses (()) and cannot be updated.
- Tuples can be thought of as read-only lists.
- For example – Tuples are fixed size in nature whereas lists are dynamic. In other words, a tuple is immutable whereas a list is mutable. You can't add elements to a tuple. Tuples have no append or extend method. You can't remove elements from a tuple. Tuples have no remove or pop method.

Python Dictionary:

- Python's dictionaries are kind of hash table type. They work like associative arrays

or hashes found in Perl and consist of key-value pairs. A dictionary key can be almost any Python type, but are usually numbers or strings. Values, on the other hand, can be any arbitrary Python object.

- Dictionaries are enclosed by curly braces ({ }) and values can be assigned and accessed using square braces ([]).
- You can use numbers to "index" into a list, meaning you can use numbers to find out what's in lists. You should know this about lists by now, but make sure you understand that you can only use numbers to get items out of a list.
- What a dict does is let you use anything, not just numbers. Yes, a dict associates one thing to another, no matter what it is.

PYTHON FUNCTION:

Defining a Function:

You can define functions to provide the required functionality. Here are simple rules to define a function in Python. Function blocks begin with the keyword `def` followed by the function name and parentheses (i.e.()).

Any input parameters or arguments should be placed within these parentheses. You can also define parameters inside these parentheses

The code block within every function starts with a colon (:) and is indented. The statement `return [expression]` exits a function, optionally passing back an expression to the caller. A return statement with no arguments is the same as `return None`.

Calling a Function:

Defining a function only gives it a name, specifies the parameters that are to be included in the function and structures the blocks of code. Once the basic structure of a function is finalized, you can execute it by calling it from another function or directly from the Python prompt.

PYTHON USING OOP's CONCEPTS:

Class:

- Class: A user-defined prototype for an object that defines a set of attributes that characterize any object of the class. The attributes are data members (class variables and instance variables) and methods, accessed via dot notation.
- Class variable: A variable that is shared by all instances of a class. Class variables are defined within a class but outside any of the class's methods. Class variables are not used as frequently as instance variables are.
- Data member: A class variable or instance variable that holds data associated with a class and its objects.
- Instance variable: A variable that is defined inside a method and belongs only to the current instance of a class.
- Defining a Class:
 - We define a class in a very similar way how we define a function.
 - Just like a function ,we use parentheses and a colon after the class name(i.e. (:)) when we define a class. Similarly, the body of our class is

indented like a functions body is.

```
def my_function():  
    # the details of the  
    # function go here
```

```
class MyClass():  
    # the details of the  
    # class go here
```

Figure 7 : Defining a Class

__init__ method in Class:

- The init method — also called a constructor — is a special method that runs when an instance is created so we can perform any tasks to set up the instance.
- The init method has a special name that starts and ends with two underscores: `init ()`.

CHAPTER 3

CASE STUDY

PROBLEM STATEMENT:

This is a difficult regression task, where the aim is to predict the burned area of forest fires, in the northeast region of Portugal.

DATA SET:

The given data set consists of the following parameters:

1. X - x-axis spatial coordinate within the Montesinho park map: 1 to 9
2. Y - y-axis spatial coordinate within the Montesinho park map: 2 to 9
3. month - month of the year: 'jan' to 'dec'
4. day - day of the week: 'mon' to 'sun'
5. FFMC - FFMC index from the FWI system: 18.7 to 96.20
6. DMC - DMC index from the FWI system: 1.1 to 291.3
7. DC - DC index from the FWI system: 7.9 to 860.6
8. ISI - ISI index from the FWI system: 0.0 to 56.10
9. temp - temperature in Celsius degrees: 2.2 to 33.30
10. RH - relative humidity in %: 15.0 to 100
11. wind - wind speed in km/h: 0.40 to 9.40
12. rain - outside rain in mm/m2 : 0.0 to 6.4
13. area - the burned area of the forest (in ha): 0.00 to 1090.84

OBJECTIVE OF THE CASE STUDY:

This is a regression problem with clear outliers which cannot be predicted using any reasonable method. A comparison of the three methods has been done :

- (a) Random Forest Regressor,
- (b) Neural Network,
- (c) Linear Regression

The output 'area' was first transformed with a $\ln(x+1)$ function.

One regression metric was measured: RMSE and r^2 score is obtained. An analysis to the regression error curve(REC) shows that the RFR model predicts more examples within a lower admitted error. In effect, the RFR model predicts better small fires and r^2 score is obtained by using Linear Regression.

CHAPTER 4

MODEL BUILDING

PREPROCESSING OF THE DATA:

Preprocessing of the data actually involves the following steps:

GETTING THE DATASET:

we can get the data from client. we can get the data from database.

IMPORTING THE LIBRARIES:

We have to import the libraries as per the requirement of the algorithm.

Importing required packages

```
[ ] import numpy as np
import pandas as pd
import seaborn as sns
import sklearn
import matplotlib.pyplot as plt
%matplotlib inline

➤ /usr/local/lib/python3.6/dist-packages/statsmodels/tools
import pandas.util.testing as tm
```


Figure 8 : Importing Libraries

IMPORTING THE DATA-SET:

Pandas in python provide an interesting method `read_csv()`. The `read_csv` function reads the entire dataset from a comma separated values file and we can assign it to a DataFrame to which all the operations can be performed. It helps us to access each and every row as well as columns and each and every value can be access using the dataframe. Any missing value or NaN value have to be cleaned.

Reading the dataset

```
[ ] data = pd.read_csv("/content/drive/My Drive/Summer Internship Project /Predict the ForestFires/forestfires.csv")
data
```




	X	Y	month	day	FFMC	DMC	DC	ISI	temp	RH	wind	rain	area
0	7	5	mar	fri	86.2	26.2	94.3	5.1	8.2	51	6.7	0.0	0.00
1	7	4	oct	tue	90.6	35.4	669.1	6.7	18.0	33	0.9	0.0	0.00
2	7	4	oct	sat	90.6	43.7	686.9	6.7	14.6	33	1.3	0.0	0.00
3	8	6	mar	fri	91.7	33.3	77.5	9.0	8.3	97	4.0	0.2	0.00
4	8	6	mar	sun	89.3	51.3	102.2	9.6	11.4	99	1.8	0.0	0.00
...
512	4	3	aug	sun	81.6	56.7	665.6	1.9	27.8	32	2.7	0.0	6.44
513	2	4	aug	sun	81.6	56.7	665.6	1.9	21.9	71	5.8	0.0	54.29
514	7	4	aug	sun	81.6	56.7	665.6	1.9	21.2	70	6.7	0.0	11.16
515	1	4	aug	sat	94.4	146.0	614.7	11.3	25.6	42	4.0	0.0	0.00
516	6	3	nov	tue	79.5	3.0	106.7	1.1	11.8	31	4.5	0.0	0.00

Figure 9: Reading the dataset

HANDLING MISSING VALUES:

Missing values

```
[ ] data.isna().sum()    ### checking missing values
```



```
X      0
Y      0
month  0
day    0
FFMC   0
DMC    0
DC     0
ISI    0
temp   0
RH     0
wind   0
rain   0
area   0
dtype: int64
```

Figure 10: Missing values



OBSERVATION:



As we can see there are no missing values in the given dataset of forest fires



DATA VISUALIZATION:


```
[ ] data['Log-area']=np.log10(data['area']+1)
```

```
[ ] for i in data.describe().columns[:-2]:  
    data.plot.scatter(i,'Log-area',grid=True)
```












Figure 11: scatterplots and distributions of numerical features to see how they may affect the output 'area'

```
[ ] data.boxplot(column='Log-area',by='day')
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f2c9ab765f8>
```

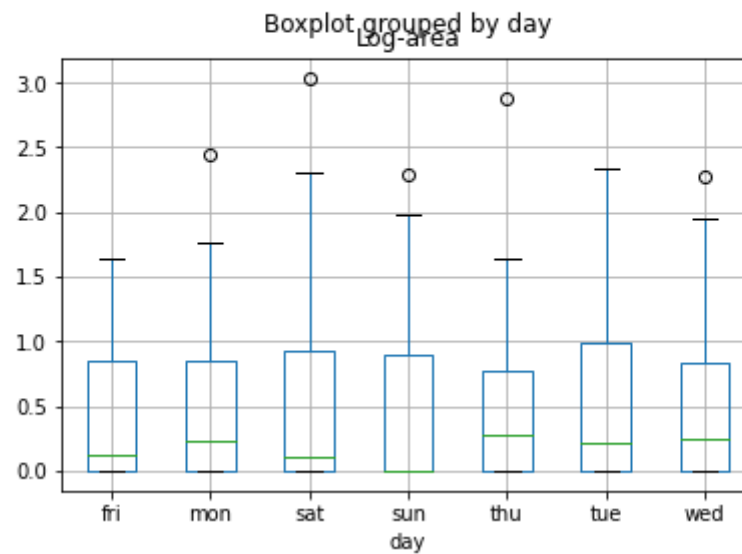


Figure 12: Boxplot of how categorical column day affect the outcome

```
[ ] data.boxplot(column='Log-area',by='month')
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f2c9a4df278>
```




Figure 13:Boxplot of how categorical column month affect the outcome

CATEGORICAL DATA:

- Machine Learning models are based on equations, we need to replace the text by numbers. So that we can include the numbers in the equations.
- Categorical Variables are of two types: Nominal and Ordinal
- **Nominal:** The categories do not have any numeric ordering in between them. They don't have any ordered relationship between each of them. Examples: Male or Female, any colour
- **Ordinal:** The categories have a numerical ordering in between them. Example: Graduate is less than Post Graduate, Post Graduate is less than Ph.D. customer satisfaction survey, high low medium
- Categorical data can be handled by using dummy variables, which are also called as indicator variables.
- Handling categorical data using dummies:

In pandas library we have a method called `get_dummies()` which creates dummy variables for those categorical data in the form of 0's and 1's.

Once these dummies got created we have to concat this dummy set to our dataframe or we can add that dummy set to the dataframe.

```
[ ] data['month']
```

```
0    mar
1    oct
2    oct
3    mar
4    mar
...
512  aug
513  aug
514  aug
515  aug
516  nov
Name: month, Length: 517, dtype: object
```

Figure 12: Categorical data-column 'month'

```
## Covert the categorical data into numerical columns using get_dummies
dummy_set = pd.get_dummies(data.month)
dummy_set
```

	apr	aug	dec	feb	jan	jul	jun	mar	may	nov	oct	sep
0	0	0	0	0	0	0	0	1	0	0	0	0
1	0	0	0	0	0	0	0	0	0	0	1	0
2	0	0	0	0	0	0	0	0	0	0	1	0
3	0	0	0	0	0	0	0	1	0	0	0	0
4	0	0	0	0	0	0	0	1	0	0	0	0
...
512	0	1	0	0	0	0	0	0	0	0	0	0
513	0	1	0	0	0	0	0	0	0	0	0	0
514	0	1	0	0	0	0	0	0	0	0	0	0
515	0	1	0	0	0	0	0	0	0	0	0	0
516	0	0	0	0	0	0	0	0	0	1	0	0

517 rows x 12 columns

Figure 13: dummy set for column 'month'

```
data['day']
```

```
0    fri
1    tue
2    sat
3    fri
4    sun
...
512  sun
513  sun
514  sun
515  sat
516  tue
Name: day, Length: 517, dtype: object
```

Figure 14: Categorical column-'day'

```
dummy_set1 = pd.get_dummies(data.day)
dummy_set1
```

	fri	mon	sat	sun	thu	tue	wed
0	1	0	0	0	0	0	0
1	0	0	0	0	0	1	0
2	0	0	1	0	0	0	0
3	1	0	0	0	0	0	0
4	0	0	0	1	0	0	0
...
512	0	0	0	1	0	0	0
513	0	0	0	1	0	0	0
514	0	0	0	1	0	0	0
515	0	0	1	0	0	0	0
516	0	0	0	0	0	1	0

517 rows x 7 columns

Figure 15: dummy set for column 'day'

```
# Concatenating the original dataframe and dataframe of dummy columns
merged_data = pd.concat([data, dummy_set, dummy_set1 ], axis=1)
merged_data
```

	X	Y	month	day	FFMC	DMC	DC	ISI	temp	RH	wind	rain	area	Log-area	apr	aug	dec	feb	jan	jul	jun	mar	may	nov	oct	sep	fri	mon	sat	si
0	7	5	mar	fri	86.2	26.2	94.3	5.1	8.2	51	6.7	0.0	0.00	0.000000	0	0	0	0	0	0	0	1	0	0	0	0	1	0	0	
1	7	4	oct	tue	90.6	35.4	669.1	6.7	18.0	33	0.9	0.0	0.00	0.000000	0	0	0	0	0	0	0	0	0	0	1	0	0	0	0	
2	7	4	oct	sat	90.6	43.7	686.9	6.7	14.6	33	1.3	0.0	0.00	0.000000	0	0	0	0	0	0	0	0	0	0	1	0	0	0	0	1
3	8	6	mar	fri	91.7	33.3	77.5	9.0	8.3	97	4.0	0.2	0.00	0.000000	0	0	0	0	0	0	0	1	0	0	0	0	1	0	0	
4	8	6	mar	sun	89.3	51.3	102.2	9.6	11.4	99	1.8	0.0	0.00	0.000000	0	0	0	0	0	0	0	1	0	0	0	0	0	0	0	0
...
512	4	3	aug	sun	81.6	56.7	665.6	1.9	27.8	32	2.7	0.0	6.44	0.871573	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0
513	2	4	aug	sun	81.6	56.7	665.6	1.9	21.9	71	5.8	0.0	54.29	1.742647	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0
514	7	4	aug	sun	81.6	56.7	665.6	1.9	21.2	70	6.7	0.0	11.16	1.084934	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0
515	1	4	aug	sat	94.4	146.0	614.7	11.3	25.6	42	4.0	0.0	0.00	0.000000	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	1
516	6	3	nov	tue	79.5	3.0	106.7	1.1	11.8	31	4.5	0.0	0.00	0.000000	0	0	0	0	0	0	0	0	0	1	0	0	0	0	0	0

Figure 16: Concatenating dummy sets to dataframe

- Getting dummies using label encoder from scikit learn package

We have a method called label encoder in scikit learn package .we need to import the label encoder method from scikitlearn package and after that we have to fit and transform the data frame to make the categorical data into dummies.

If we use this method to get dummies then in place of categorical data we get the numerical values (0,1,2....)


```
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import OneHotEncoder
from sklearn.preprocessing import LabelEncoder
```

```
enc = LabelEncoder()
enc.fit(data['month'])    ## month column
```

```
LabelEncoder()
```

Figure 17: importing label encoder and one hot encoder

```
enc.classes_
array(['apr', 'aug', 'dec', 'feb', 'jan', 'jul', 'jun', 'mar', 'may',
      'nov', 'oct', 'sep'], dtype=object)
```

```
data['month_encoded']=enc.transform(data['month'])
data.head()
```

	X	Y	month	day	FFMC	DMC	DC	ISI	temp	RH	wind	rain	area	Log-area	month_encoded
0	7	5	mar	fri	86.2	26.2	94.3	5.1	8.2	51	6.7	0.0	0.0	0.0	7
1	7	4	oct	tue	90.6	35.4	669.1	6.7	18.0	33	0.9	0.0	0.0	0.0	10
2	7	4	oct	sat	90.6	43.7	686.9	6.7	14.6	33	1.3	0.0	0.0	0.0	10
3	8	6	mar	fri	91.7	33.3	77.5	9.0	8.3	97	4.0	0.2	0.0	0.0	7
4	8	6	mar	sun	89.3	51.3	102.2	9.6	11.4	99	1.8	0.0	0.0	0.0	7

Figure 18: Handling categorical data of column month

```
enc.classes_
```

```
array(['fri', 'mon', 'sat', 'sun', 'thu', 'tue', 'wed'], dtype=object)
```

```
data['day_encoded']=enc.transform(data['day'])  
data.head(10)
```

	X	Y	month	day	FFMC	DMC	DC	ISI	temp	RH	wind	rain	area	Log-area	month_encoded	day_encoded
0	7	5	mar	fri	86.2	26.2	94.3	5.1	8.2	51	6.7	0.0	0.0	0.0	7	0
1	7	4	oct	tue	90.6	35.4	669.1	6.7	18.0	33	0.9	0.0	0.0	0.0	10	5
2	7	4	oct	sat	90.6	43.7	686.9	6.7	14.6	33	1.3	0.0	0.0	0.0	10	2
3	8	6	mar	fri	91.7	33.3	77.5	9.0	8.3	97	4.0	0.2	0.0	0.0	7	0
4	8	6	mar	sun	89.3	51.3	102.2	9.6	11.4	99	1.8	0.0	0.0	0.0	7	3
5	8	6	aug	sun	92.3	85.3	488.0	14.7	22.2	29	5.4	0.0	0.0	0.0	1	3
6	8	6	aug	mon	92.3	88.9	495.6	8.5	24.1	27	3.1	0.0	0.0	0.0	1	1
7	8	6	aug	mon	91.5	145.4	608.2	10.7	8.0	86	2.2	0.0	0.0	0.0	1	1
8	8	6	sep	tue	91.0	129.5	692.6	7.0	13.1	63	5.4	0.0	0.0	0.0	11	5
9	7	5	sep	sat	92.5	88.0	698.6	7.1	22.8	40	4.0	0.0	0.0	0.0	11	2

Figure 19: Handling categorical data of column day

TRAINING THE MODEL:

- Splitting the data : after the preprocessing is done then the data is split into train and test sets
- In Machine Learning in order to access the performance of the classifier. You train the classifier using 'training set' and then test the performance of your classifier on unseen 'test set'. An important point to note is that during training the classifier only uses the training set . The test set must not be used during training the classifier. The test set will only be available during testing the classifier.
- training set - a subset to train a model.(Model learns patterns between Input and Output)
- test set - a subset to test the trained model.(To test whether the model has correctly learnt)
- The amount or percentage of Splitting can be taken as specified (i.e. train data = 75% , test data =25% or train data = 80% , test data= 20%)
- First we need to identify the input and output variables and we need to separate the input set and output set
- In scikit learn library we have a package called model_selection in which train_test_split method is available .we need to import this method
- This method splits the input and output data to train and test based on the percentage specified by the user and assigns them to four different variables(we need to mention the variables)

```
# Preparing Training and Testing Data
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.4,
                                                    random_state= 156)

print(X_train.shape)
print(X_test.shape)
print(y_train.shape)
print(y_test.shape)
X_train
```

(310, 12)
(207, 12)
(310,)
(207,)

Figure 20: Importing train_test_split

Building the model (using splitting):

- First we have to retrieve the input and output sets from the given dataset

```
X = data.drop(['area', 'Log-area', 'month', 'day'], axis=1)
X.head()
```

	X	Y	FFMC	DMC	DC	ISI	temp	RH	wind	rain	month_encoded	day_encoded
0	7	5	86.2	26.2	94.3	5.1	8.2	51	6.7	0.0	7	0
1	7	4	90.6	35.4	669.1	6.7	18.0	33	0.9	0.0	10	5
2	7	4	90.6	43.7	686.9	6.7	14.6	33	1.3	0.0	10	2
3	8	6	91.7	33.3	77.5	9.0	8.3	97	4.0	0.2	7	0
4	8	6	89.3	51.3	102.2	9.6	11.4	99	1.8	0.0	7	3

Figure 21: Retrieving the input columns

```

y = data['Log-area']
y.head()

0    0.0
1    0.0
2    0.0
3    0.0
4    0.0
Name: Log-area, dtype: float64

```

Figure 22: Retrieving output column

MODEL BUILDING:

```

def rec(m,n,tol):
    if type(m)!='numpy.ndarray':
        m=np.array(m)
    if type(n)!='numpy.ndarray':
        n=np.array(n)
    l=m.size
    percent = 0
    for i in range(l):
        if np.abs(10**m[i]-10**n[i])<=tol:
            percent+=1
    return 100*(percent/l)

tol_max=20

```

Figure 23: Defining Regression Error Characteristic (REC)

Regression Error Characteristic (REC) estimation:

Receiver Operating Characteristic (ROC) curves provide a powerful tool for visualizing and comparing classification results. Regression Error Characteristic (REC) curves generalize ROC curves to regression. REC curves plot the error tolerance on the versus the percentage of points predicted within the tolerance on the . The resulting curve estimates the cumulative distribution function of the error. The REC curve visually presents commonly-used statistics. The area-over-the-curve (AOC) is a biased estimate of the expected error.

The value can be estimated using the ratio of the AOC for a given model to the AOC for the nul-model. Users can quickly assess the relative merits of many regression functions by examining the relative position of their REC curves. The shape of the curve reveals additional information that can be used to guide modeling.

RANDOM FOREST REGRESSOR

```
[ ] from sklearn.ensemble import RandomForestRegressor
    from sklearn.model_selection import GridSearchCV
```

Figure 24: Importing random forest regressor and gridsearchCV

Gridsearch

Finding the right parameters for machine learning models is a tricky task. But luckily, Scikit-learn has the functionality of trying a bunch of combinations and see what works best, built in with GridSearchCV. The CV stands for cross-validation.

GridSearchCV takes a dictionary that describes the parameters that should be tried and a model to train. The grid of parameters is defined as a dictionary, where the keys are the parameters and the values are the settings to be tested.

```
# Parameter grid for the Grid Search
param_grid = {'C': [0.01,0.1,1, 10], 'epsilon': [10,1,0.1,0.01,0.001,0.0001], 'kernel': ['rbf']}

param_grid = {'max_depth': [5,10,15,20,50], 'max_leaf_nodes': [2,5,10], 'min_samples_leaf': [2,5,10],
              'min_samples_split':[2,5,10]}
grid_RF = GridSearchCV(RandomForestRegressor(),param_grid,refit=True,verbose=0,cv=5)
grid_RF.fit(X_train,y_train)

GridSearchCV(cv=5, error_score=nan,
             estimator=RandomForestRegressor(bootstrap=True, ccp_alpha=0.0,
             criterion='mse', max_depth=None,
             max_features='auto',
             max_leaf_nodes=None,
             max_samples=None,
             min_impurity_decrease=0.0,
             min_impurity_split=None,
             min_samples_leaf=1,
             min_samples_split=2,
             min_weight_fraction_leaf=0.0,
             n_estimators=100, n_jobs=None,
             oob_score=False, random_state=None,
```

Figure 25: parameter grid for grid search and rfr

```
print("Best parameters obtained by Grid Search:",grid_RF.best_params_)
```

Best parameters obtained by Grid Search: {'max_depth': 20, 'max_leaf_nodes': 2, 'min_samples_leaf': 5, 'min_samples_split': 2}

Figure 26: best parameters obtained by grid search

```
a=grid_RF.predict(X_test)
rmse_rf=np.sqrt(np.mean((y_test-a)**2))
print("RMSE for Random Forest:",rmse_rf)
```

RMSE for Random Forest: 0.6282077936297361

Figure 27: RMSE for random forest

```
plt.xlabel("Actual area burned")
plt.ylabel("Error")
plt.grid(True)
plt.scatter(10**(y_test),10**(a)-10**(y_test))
```

<matplotlib.collections.PathCollection at 0x7f5a3be935f8>




Figure 28: scatter plot
x-axis: actual area burned
y-axis: error

```
plt.title("Histogram of prediction errors\n",fontsize=18)
plt.xlabel("Prediction error ($ha$)",fontsize=14)
plt.grid(True)
plt.hist(10**(a.reshape(a.size,))-10**(y_test),bins=50)
```

```
(array([ 1.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,
        0.,  0.,  1.,  0.,  0.,  0.,  0.,  1.,  0.,  0.,  0.,
        1.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,  1.,  0.,  1.,
        0.,  1.,  0.,  0.,  2.,  1.,  0.,  1.,  1.,  4.,  4.,
        3.,  4.,  2., 15., 27., 136.]),
 array([-276.62645053, -271.02039101, -265.41433149, -259.80827198,
        -254.20221246, -248.59615295, -242.99009343, -237.38403392,
        -231.7779744 , -226.17191489, -220.56585537, -214.95979586,
        -209.35373634, -203.74767683, -198.14161731, -192.5355578 ,
        -186.92949828, -181.32343877, -175.71737925, -170.11131974,
        -164.50526022, -158.89920071, -153.29314119, -147.68708168,
        -142.08102216, -136.47496265, -130.86890313, -125.26284361,
        -119.6567841 , -114.05072458, -108.44466507, -102.83860555,
        -97.23254604, -91.62648652, -86.02042701, -80.41436749,
        -74.80830798, -69.20224846, -63.59618895, -57.99012943,
        -52.38406992, -46.7780104 , -41.17195089, -35.56589137,
        -29.95983186, -24.35377234, -18.74771283, -13.14165331,
        -7.5355938 , -1.92953428,  3.67652523]),
 <a list of 50 Patch objects>)
```

Histogram of prediction errors




Figure 29: Histogram of prediction errors of rfr


```

rec_RF=[]
for i in range(tol_max):
    rec_RF.append(rec(a,y_test,i))

plt.figure(figsize=(5,5))
plt.title("REC curve for the Random Forest\n",fontsize=15)
plt.xlabel("Absolute error (tolerance) in prediction ($ha$)")
plt.ylabel("Percentage of correct prediction")
plt.xticks([i for i in range(0,tol_max+1,5)])
plt.ylim(-10,100)
plt.yticks([i*20 for i in range(6)])
plt.grid(True)
plt.plot(range(tol_max),rec_RF)

```

[<matplotlib.lines.Line2D at 0x7f5a3bfc5ac8>]

REC curve for the Random Forest




Figure 30: REC curve for rfr

NEURAL NETWORK

```
[ ] from keras.models import Sequential
    import keras.optimizers as opti
    from keras.layers import Dense, Activation, Dropout
```

☞ Using TensorFlow backend.

Figure 31: Importing NN packages

```
data=X_train
target = y_train
model.fit(data, target, epochs=100, batch_size=10, verbose=0)
```

<keras.callbacks.callbacks.History at 0x7f59f93a9da0>

Figure 32: dividing data and target

Prediction and RMSE

```
[ ] a=model.predict(X_test)
    print("RMSE for NN:", np.sqrt(np.mean((y_test-a.reshape(a.size,))**2)))
```

☞ RMSE for NN: 0.6243784239764546

Figure 33: RMSE for NN

The RMSE of rfr is 0.628

The RMSE of NN is 0.624

Random Forest Regressor algorithm in which GridsearchCV is used is the best one to predict.

```
plt.xlabel("Actual area burned")
plt.ylabel("Error")
plt.grid(True)
plt.scatter(10**(y_test), 10**(a.reshape(a.size,))-10**(y_test))
```

<matplotlib.collections.PathCollection at 0x7f59f1f4ec50>




Figure 34: scatterplot for NN
x-axis: actual area burned
y-axis: error

```
plt.title("Histogram of prediction errors\n",fontsize=16)
plt.xlabel("Prediction error ($ha$)",fontsize=14)
plt.grid(True)
plt.hist(10**(a.reshape(a.size,))-10**(y_test),bins=50)

(array([ 1.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,
        0.,  0.,  1.,  0.,  0.,  0.,  0.,  1.,  0.,  0.,  0.,
        1.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,  1.,  1.,
        0.,  0.,  1.,  0.,  2.,  1.,  0.,  1.,  1.,  3.,  4.,
        4.,  3.,  1., 11., 23., 146.]),
 array([-276.19956568, -270.62809942, -265.05663317, -259.48516691,
        -253.91370065, -248.3422344 , -242.77076814, -237.19930189,
        -231.62783563, -226.05636938, -220.48490312, -214.91343686,
        -209.34197061, -203.77050435, -198.1990381 , -192.62757184,
        -187.05610559, -181.48463933, -175.91317307, -170.34170682,
        -164.77024056, -159.19877431, -153.62730805, -148.0558418 ,
        -142.48437554, -136.91290928, -131.34144303, -125.76997677,
        -120.19851052, -114.62704426, -109.055578 , -103.48411175,
        -97.91264549, -92.34117924, -86.76971298, -81.19824673,
        -75.62678047, -70.05531421, -64.48384796, -58.9123817 ,
        -53.34091545, -47.76944919, -42.19798294, -36.62651668,
        -31.05505042, -25.48358417, -19.91211791, -14.34065166,
        -8.7691854 , -3.19771915,  2.37374711]),
 <a list of 50 Patch objects>)
```

Histogram of prediction errors




Figure 35: Histogram of prediction error of NN

```

rec_NN=[]
for i in range(tol_max):
    rec_NN.append(rec(a,y_test,i))
plt.figure(figsize=(5,5))
plt.title("REC curve for Neural Network\n",fontsize=16)
plt.xlabel("Absolute error (tolerance) in prediction ($ha$)")
plt.ylabel("Percentage of correct prediction")
plt.xticks([i for i in range(0,tol_max+1,5)])
plt.ylim(-10,100)
plt.yticks([i*20 for i in range(6)])
plt.grid(True)
plt.plot(range(tol_max),rec_NN)
[<matplotlib.lines.Line2D at 0x7f59f1dbeda0>]

```

REC curve for Neural Network




Figure 36: REC curve for NN

Relative performance of Random Forest Regressor and Neural Networks(REC Curves)

```
[ ] plt.figure(figsize=(10,8))
plt.title("REC curve for RFR and NN\n",fontsize=20)
plt.xlabel("Absolute error (tolerance) in prediction ($ha$)",fontsize=15)
plt.ylabel("Percentage of correct prediction",fontsize=15)
plt.xticks([i for i in range(0,tol_max+1,1)],fontsize=13)
plt.ylim(-10,100)
plt.xlim(-2,tol_max)
plt.yticks([i*20 for i in range(6)],fontsize=18)
plt.grid(True)
plt.plot(range(tol_max),rec_RF,'--',lw=3)
plt.plot(range(tol_max),rec_NN,'-',lw=3)
plt.legend(['Random Forest','NN'],fontsize=13)
```

REC curve for RFR and NN




Figure 37: Relative performance of Random forest regressor and Neural Network(ROC curves)

LINEAR REGRESSION

```
[ ] from sklearn.linear_model import LinearRegression
    lm = LinearRegression()
    lm.fit(X, y)
```

➞ `LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)`

Figure 38: Importing Linear Regression package

- Import linear regression method which is available in linear_model package from scikit learn library
- Once the model is built we need to check for accuracy.
- This can be done using predict method which is used to predict the output.

```

y_pred = lm.predict(X)
y_pred
array([ 0.50557826,  0.46001527,  0.44558651,  0.3219269 ,  0.27158206,
        0.52957826,  0.50197583,  0.3381386 ,  0.65246761,  0.58976145,
        0.66906396,  0.42325735,  0.4599566 ,  0.51761182,  0.69214357,
        0.5693624 ,  0.50438361,  0.5706437 ,  0.49170854,  0.39055137,
        0.57036316,  0.4621742 ,  0.05646362,  0.57051842,  0.49560189,
        0.51069148,  0.6034448 ,  0.4809869 ,  0.60100744,  0.59993799,
        0.48078957,  0.56598604,  0.4545284 ,  0.5313777 ,  0.49057701,
        0.54941664,  0.58943755,  0.48169912,  0.55341174,  0.40066982,
        0.47231497,  0.51705945,  0.41350284,  0.4625678 ,  0.5109871 ,
        0.55523917,  0.48235461,  0.43256456,  0.33573103,  0.46691785,
        0.52076422,  0.41439744,  0.50190238,  0.50190238,  0.45945572,
        0.5440068 ,  0.52566861,  0.39498084,  0.37658534,  0.16059889,
        0.4139847 ,  0.37708504,  0.43706469,  0.39201038,  0.40559733,
        0.45425903,  0.45957251,  0.49357814,  0.52518984,  0.4275953 ,
        0.50807344,  0.41534892,  0.50154163,  0.51589654,  0.44102067,
        0.33348256,  0.42516267,  0.25243503,  0.20086118,  0.57455774,
        0.47722152,  0.32778845,  0.32788812,  0.28527164,  0.40598954,
        0.44814098,  0.47498143,  0.45425709,  0.51790374,  0.51285336,
        0.48256647,  0.55370347,  0.55209697,  0.5628246 ,  0.54003051,
        0.40230334,  0.26326514,  0.45328814,  0.34205822,  0.49292365,
        0.49292365,  0.45737736,  0.3662692 ,  0.54953436,  0.20784842,
        0.46025983,  0.42513699,  0.45986298,  0.49035207,  0.52486178,
        0.45812995,  0.42452221,  0.52910232,  0.43342139,  0.59317525,
        0.55289452,  0.48219788,  0.44618158,  0.4045798 ,  0.40246509,
        0.26931955,  0.46151367,  0.58066806,  0.49597874,  0.43087732,

```

Figure 39: predicting the output


```
[ ] y==y_pred
```

0	False
1	False
2	False
3	False
4	False
...	
512	False
513	False
514	False
515	False
516	False

Name: Log-area, Length: 517, dtype: bool

Figure 40: Comparing y with y_pred

```
[ ] ## r2_value--> to check the model performance
    from sklearn.metrics import r2_score
    r2_score(y, y_pred)
```

```
0.027927989936573083
```

Figure 41: r2_score to check the model performance

The r2_score obtained is 0.02

CONCLUSION:

Forest fires cause a significant environmental damage while threatening human lives. In the last two decades, a substantial effort was made to build automatic detection tools that could assist Fire Management Systems (FMS). The three major trends are the use of satellite data, infrared/smoke scanners and local sensors (e.g. meteorological). The advantage is that such data can be collected in real-time and with very low costs, when compared with the satellite and scanner approaches. Recent real-world data, from the northeast region of Portugal, was used in the experiments. The database included spatial, temporal, components from the Canadian Fire Weather Index (FWI) and four weather conditions.

This problem was modeled as a regression task, where the aim was the prediction of the burned area. The drawback is the lower predictive accuracy for large fires. To our knowledge, this is the first time the burn area is predicted using only meteorological based data and further exploratory research is required. As argued in , predicting the size of forest fires is a challenging task. To improve it, we believe that additional information is required, such as the type of vegetation and firefighting intervention (e.g. time elapsed and firefighting strategy). Since the FWI system is widely used around the world, further research is need to confirm if direct weather conditions are preferable than accumulated values, as suggested by this study.

Finally, since large fires are rare events, outlier detection techniques will also be addressed.

REFERENCES:

- [1] <https://archive.ics.uci.edu/ml/datasets/forest+fires>

- [2] https://en.wikipedia.org/wiki/Machine_learning

- [3] https://en.wikipedia.org/wiki/Deep_learning